Product Forum: “GA Readiness” for an AI-first world

Anna Pope, Kaleem Rahman

January 26, 2025

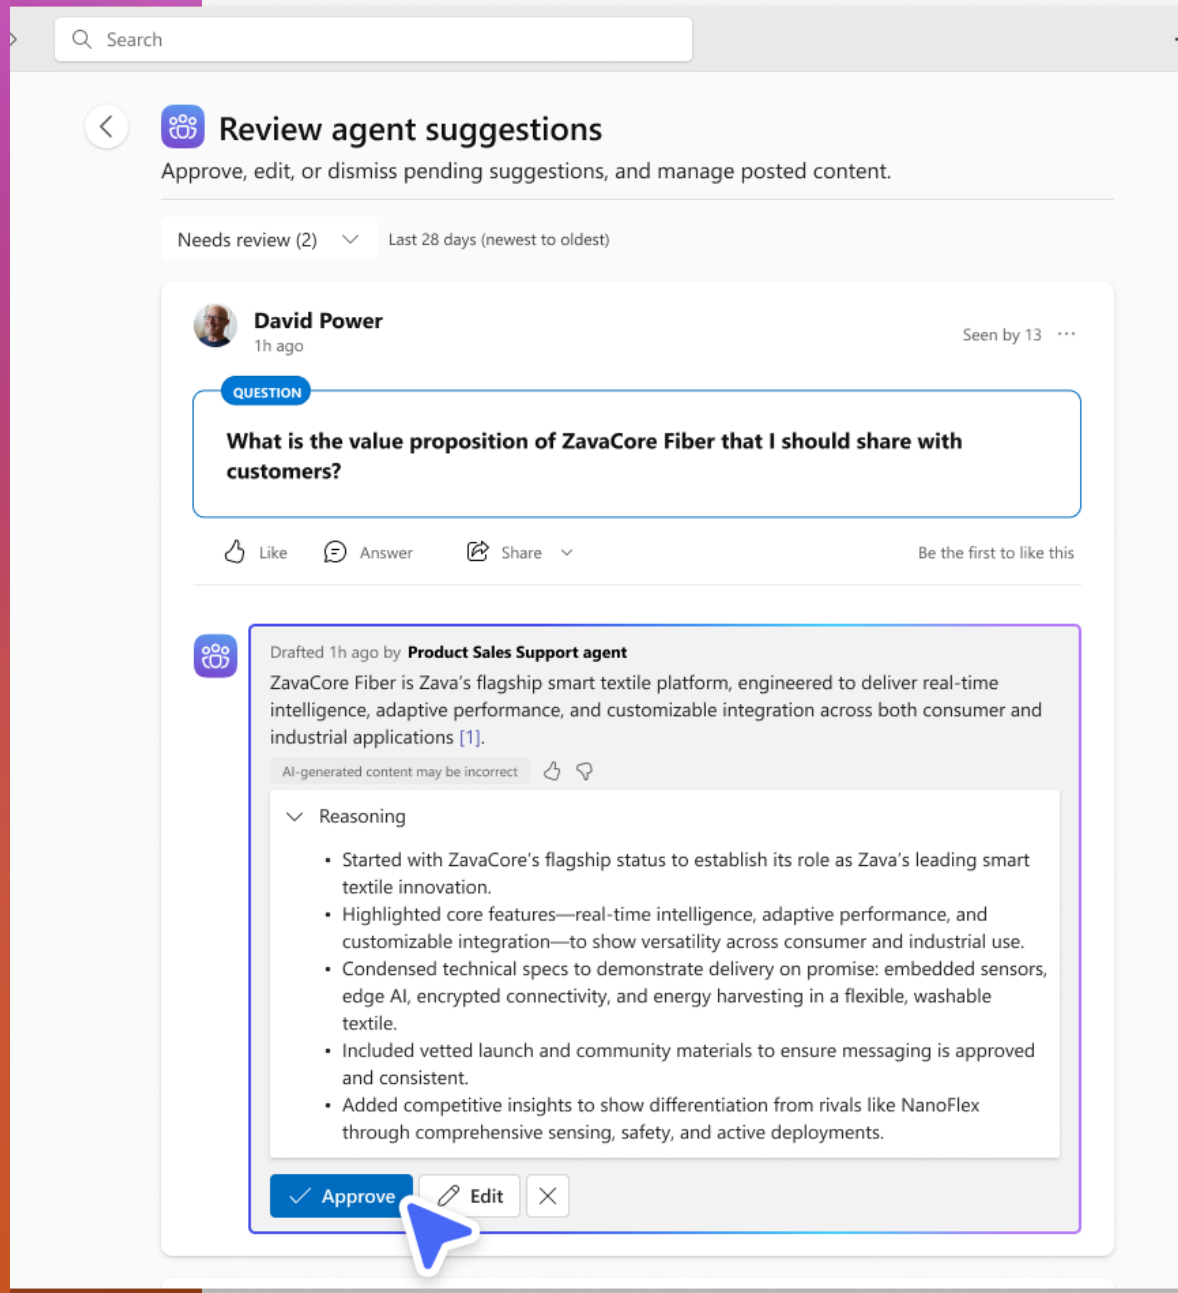

Learning 1: Accuracy alone ≠ perceived quality

- Traditionally, quality = reliability, polished UX, few bugs
- We prioritized accuracy first ("people won't trust it if wrong"); thanks to AS, this is solid
- However, preview customers complained:
  - "Answers are too long"
  - "Sounds robotic"
  - "Uses inappropriate tone"
- You need to work out what this means. For example, desirable tone =
  - Starts by acknowledging question
  - Write as "we" or "I", not a third-person expert
  - Uses contractions (e.g. shouldn't)
  - Empathetic and personable
- You've to identify the factors that are consequential to *your* scenario by evaluating agent outputs in numbers, getting customer signal and clear articulation

Learning 2: Quality systems need evaluators + datasets + calibration

- **Datasets** (representative scenarios to test), **evaluators** (identifying what "good" means) and **calibration** (tuning, thresholds, judge alignment) work together
- Two evaluator families:
  - **Reference-based** (compare to a "golden answer") and **reference-free** (using a "formula" to judge)
  - Implemented using code or an LLM-as-Judge
- **The loop**: define metric → implement as judge or code → test against human labels → calibrate
  - Calibration important to build trust and can take time
- **Deciding where to invest in is consequential**
  - Resist temptation to take on generic, out-of-the-box AI metrics like "helpfulness", "coherence", "toxicity"
  - Focus on *your* most consequential failure modes or quality dimensions

Learning 3: GA Readiness Is a Function of Content and Workflow

Content Quality: Is the output useful?
- Useful, low validation effort, grounded
- Key Metrics: Thumbs Up/Down, Posted % (keep rate)

Workflow fit: Does it reduce user effort?
- Workflow must enable evaluation, not obscure it.
- Key Metrics: Pending %, Attach rate, Retention

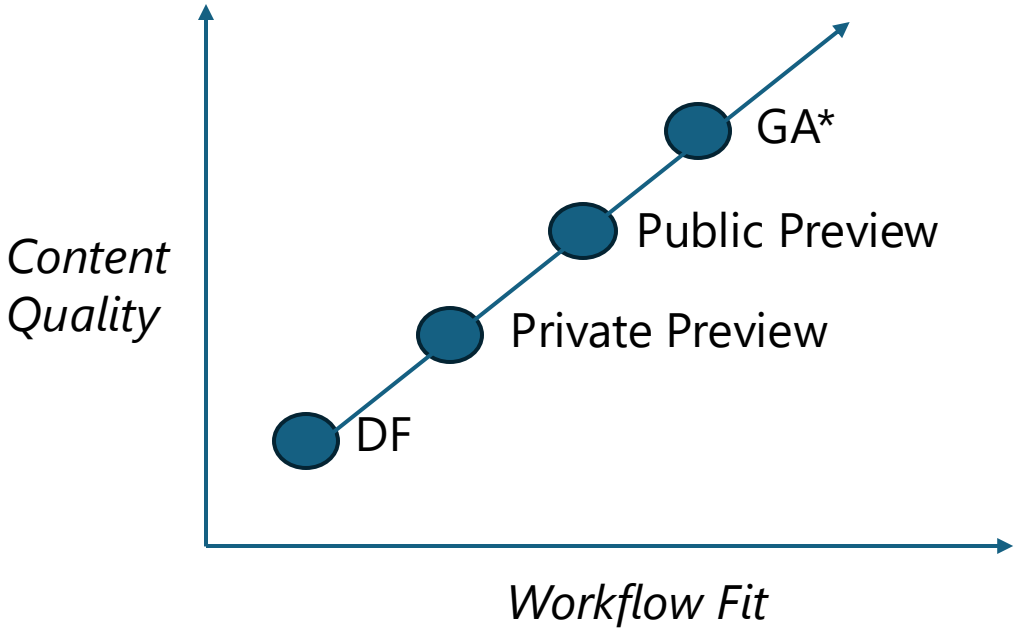

Quality has the longest lead time due to experimentation and signal generation, so focus on learning first while making UX refinements in parallel.

*GA requires ongoing monitoring and calibration

Facilitator | GA Dashboard

JTBD.	Skills GA	In R4?	Reliability (R4)	Latency P90 (R4)	Legibility/ MMLU (4-5 good)	Task Effectiveness (>4 good)	Intent Confidence (>85% good)	Fairy Dust/ Perceived Intelligence (≥4 good)	Usage (R7) (includes GA)	Try Retention (includes GA)	SAT
Capture meeting notes & action items	1. Live notetaking	Yes	99.29	NA	4.2*	NA	NA	NA	47.9k (R4: 17.4K)	34.1%	NA
Establish shared goal & agenda	2. Extract agenda from meeting invite	Yes	99.95	NA	NA	4.4**	NA	NA	31.08K	NA	84.7%
	3. Extract agenda from Loop	Yes	99.75	-	NA	4.8**	NA	NA	931		
	4. If no agenda exists, establish meeting goal from conversation	Yes	99.95	NA	-	4.45**	NA	-			
Stay on agenda, end on time	5. Mid-point/wrap up time check for meetings without agendas	Yes	99.95	NA	-	4.6**	NA	-	63.8k Meetings	NA	
	6. Agenda timer for agendas with time (proactive)	Yes	99.95	NA	NA	NA	NA	-			
	7. Track discussion for agendas without time (proactive)	Yes	99.95	NA	NA	NA	NA	-			
	8. Set & control ad-hoc timer	Yes	99.45	17.5s (4.0s)	NA	4.5*	NA	-			
	9. Overtime alert with scheduling option	Yes	99.74	NA	NA	4.3**	NA	-			
	10. Agenda Timer (w/ or wt/ time) thru chat commanding	Yes	TBD	2.1s 1	NA	4.5*	NA	-			
Get quick answers & find relevant content	11. Ask questions about meeting	Yes	99.16	19s	4.32*	NA	98%	4.17	892 (R4: 350)	14.6%	76.1%
	12. Ask questions about chat	Yes	99.54	12.8s	4.43*	NA	87%	4.11	450 (R4: 239)	5.5%	
	13. Ask questions about web	Yes	100	14.7s	4.21*	NA	96%	3.92	250 (R4: 140)	5.5%	
	14. Handle Banter and Small Talk	Yes	NA	NA	4.19*	NA	94%	3.8			
	15. Ask questions about time, agenda and people in meeting	Yes			4.19*	NA	91%	4.61	738 (R4: 360)	10.7%	
	16. Ask questions about Meeting Agent	Yes	100	9.9s	4.12*	NA	91%	4.31	337 (R4: 221)	3.5%	
	17. Soft Disengagement when skill is unsupported	Yes			4.46*	NA	94%	4.12			

Project Manager agent scorecard

Comparing week ending Oct 4 with week ending Oct 11

Top PMA tenants

1. **Accenture** | 210 WAU | 13.81% W1 retention
2. **EY** | 157 WAU | 9.55% W1 retention
3. **KPMG** | 127 WAU | 17.32% W1 retention
4. **NOVO NORDISK** | 72 WAU | 6.94% W1 retention
5. **DELOITTE** (O365D) | 88 WAU | 6.82% W1 retention
6. **SIEMENS AG** | 92 WAU | 13.04% W1 retention
7. **NESTLE** | 58 WAU | 8.62% W1 retention
8. **VISA INC.** | 54 WAU | 12.96% W1 retention
9. **NOVARTIS** PHARMA AG | 69 WAU | 13.04% W1 retention
10. **PETROBRAS** | 39 WAU | 5.13% W1 retention

In Planner	AI Adoption			AI Quality	AI Health	
Scenario *Skills*	**Usage** Current 34.6K MAU | 11.5K WAU FY25 Baseline 6.6K WAU FY26 Goal 100K+ WAU	**Retention** Current 21.34% W1 FY25 Baseline 14.4% W1 FY26 Goal 40%+ W1	**Satisfaction** Current 44% FY25 Baseline 45% FY26 Goal 70%+	**Offline Eval sets**[2] Target 4+	**Reliability** Target 99.9%	**Latency** Target is scenario dependent
Status Report *Generate a Status Report*	**8,520 MAU** +1.8% MoM, △ not stat sig WAU: 2,675 (-1.83% WoW, △ not stat sig)	**20.48% W1** -1.90% WoW, △ not stat sig Retained 558 out of 2724 With 95% CI	**61%** +0.23% MoM, △ not stat sig 474.14 per 100K tries, +1.00% MoM Likes: 99 | Dislikes: 64 With 90% CI	**4.5 LLM score** mean, with 95% CI 8 Evals: Goal alignment 4.03 | Audience fit 4.47 | Groundedness 4.7 | Insightfulness 3.9 | Consistency 4.93 | Clarity & Readability 4.97 | Fluency 4.3 | Coherence 4.8	**93.65%** last 30 days Generation reliability breakdown: • Server-side failures: 3.1% - mostly loop generation issues (recently fixed) • Client-side failures: 3.2% - roaming service failures	**191 sec** P95 Ok since report generation is async. This contain P95 of success with 177 and failure with 323.
Plan Generation *Generate a set of tasks for a goal*	**6,060 MAU** +1.7% MoM, △ not stat sig WAU: 1,893 (-2.87% WoW, △ not stat sig)	**20.5% W1** -1.4% WoW, △ not stat sig Retained 401 out of 1,949 With 95% CI	**41%** -3.3% MoM, △ not stat sig 1274 per 100K tries, +1.2% MoM Likes: 70 | Dislikes: 98 With 90% CI	**3.71 LLM score** mean, with 95% CI 2 Evals: W/ File Grounding: 3.88 | Wo/ File Grounding: 3.49 Human as a judge score: 3.78	**96.9%** last 30 days RAI err: 0.07% | 400 err: 1.69%	**60 sec** P95 Rolling out improvements.
Task Management & Q&A *Create, read, update, delete tasks + Q&A in App Chat*	**22,099 MAU** -2.6% MoM, △ not stat sig WAU: 6,469 (-2.15% WoW, △ not stat sig)	**14.81% W1** -0.18% WoW, △ not stat sig Retained 979 out of 6,609 With 95% CI	**43%** +1.24% MoM, △ not stat sig 637.39 per 100K tries, -0.10% MoM Likes: 315 | Dislikes: 426 With 95% CI	N/A This is on "legacy" App Copilot tech. We are moving to BizChat soon where we will be implementing them in SEVAL	**96.7%** last 30 days New work will be on Bizchat	**40 sec** P95 New work will be on Bizchat
Task Execution *Assign a task to PMA to create a draft*	**10,946 MAU** +2.3% MoM, △ not stat sig WAU: 3,675 (+2.68 WoW, △ not stat sig)	**23.17% W1** -1.09% WoW, △ not stat sig Retained 829 out of 3578 With 95% CI	**38%** +0.84% MoM, △ not stat sig 1022.81 per 100K tries, +9.30% MoM Likes: 129 | Dislikes: 213 With 90% CI	**3.68 LLM score** LLM-as-a-judge, grounded + ungrounded 1 Eval	**99.5%** last 30 days RAI err: 0.24% | 400 err: 1.47%	**158 sec** P95 Ok since it's async + toast notification

1. Metrics are being assessed for a Confidence Interval (CI) and changes are being assessed for stat sig, so that these numbers can be trust-worthy | 2. As we move skills to AugLoop or Sydney they will be onboarded to PRISM or SEVAL for online Evals. Currently scenarios that run on MARS only have offline evals

Community Agent: Quality-based GA readiness

We are quality-driven: GA is primarily gated by measured signals and confidence thresholds

Metric	Why Important	Current	Target
Answer posting rate % of drafted answers that ultimately get posted	Indicates agent is consistently producing publish-worthy content	**23.2%** (+1.8 pts MoM)	**40%**
Answer content quality Measured through offline (and online) evals	Measures presence of failure modes affecting quality as indicated by customers and observed by squad	Fact retrieval: **69%** Fact representation: **77%** Cognitive burden: **49%** Further reading links: **TBD** Style & formatting: **TBD** Conciseness: **TBD**	**80%***
Reviewer retention Engaged, acting on drafts	Measures whether admins and reviewers are seeing value, forming habits	**50.9%** (-0.9 pts MoM)	**80%**
👍🏻 **Thumbs up rate** % of thumbs up/down reactions that are positive	Proxy for perceived usefulness by readers	**50.1%** (+7.1 pts MoM)	**70%****

**Aspirational target to pass 4/5 cases; calibrations needed for some; there will be exceptions like fact retrieval as we adjust conciseness*
***Will materialize as we productionize evaluators*
****Standard across other ODSP agents*

Discussion

- Which GA readiness signals should persist from traditional software, and which should change?
- Is the classic “Preview → GA” lifecycle sufficient for AI, or do we need a different lifecycle vocabulary and gates?
- Should we enable more customer-visible evaluation data during previews?

Goal 2: Improve answer acceptance: ‘workflow’ enhancement

Increase reviewer action on agent drafts (approve/edit/dismiss) by making review work easier to notice, faster to complete, and lower-noise.

Experiment/Demo*	Intended Learning
In-feed indicator for pending suggestions 🧪 Surface review work in-context on questions with pending drafts	Whether in-context surfacing increase the % of reviewers who notice a draft exists and take an action
Queue count + dynamic burndown Make progress visible so the queue feels finite; celebrative UI	Whether visible progress reduces the “endless queue” effect and increase persistence / actions per visit
Reviewers can transparently see what questions were skipped or dismissed	—
Auto-remove stale drafts after best/verified Remove redundant proposals when a best/verified answer is marked	Whether removing redundant/stale items makes the queue feel manageable and drives higher action rate
Rank queue items by reviewer relevance 🧪 Prioritize what’s most relevant/important first	Whether ordering review queue by personal relevance increase actions
Lightweight sharing / routing proposals Share with another reviewer so the “right person” acts	Whether explicit routings reduce “not empowered” stalls and increases total actions
One-click edits to reduce reviewer effort** 🧪 E.g. “make concise”, custom instructions	Whether reducing rewrite effort increases approvals and reduces time-to-approve

🧪 *Launch as A/B test*

**Also in our backlog: out-of-UI nudges (Teams, Outlook), motivation & recognition, improvements that make content easier to trust and validate*

***Evaluating offline “pre-experiment” results for conciseness variants to determine if this would be meaningful to pursue*

Build your own card: Takeaways

- “Think about”
  - A: What signals indicate your solution is assisting in user goals?
  - B: How will you know users are satisfied with your solution?
  - C: What evaluators might you use to indicate the quality of your outputs?
- Come up with your top 1-3 things to keep an eye out for readiness for your AI experience. If you don’t work on an AI experience, think of your most used AI experience e.g. Claude Code, Copilot, ChatGPT, etc.