Community Agent: Cycle 6 LT Sponsor Review

January 21, 2025

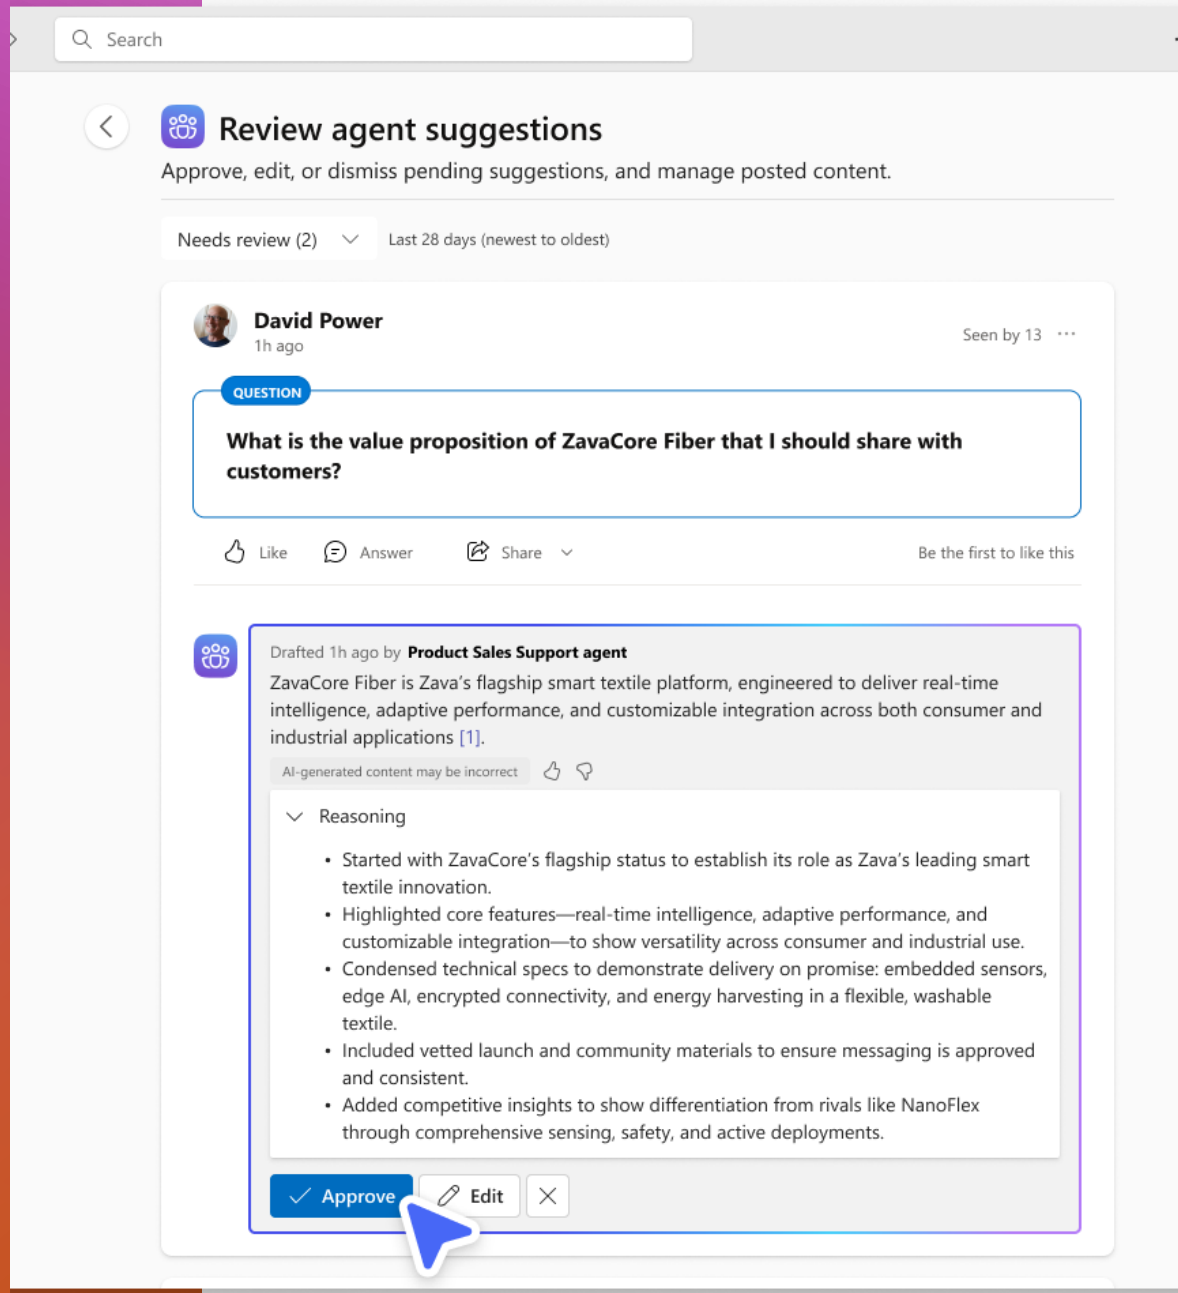

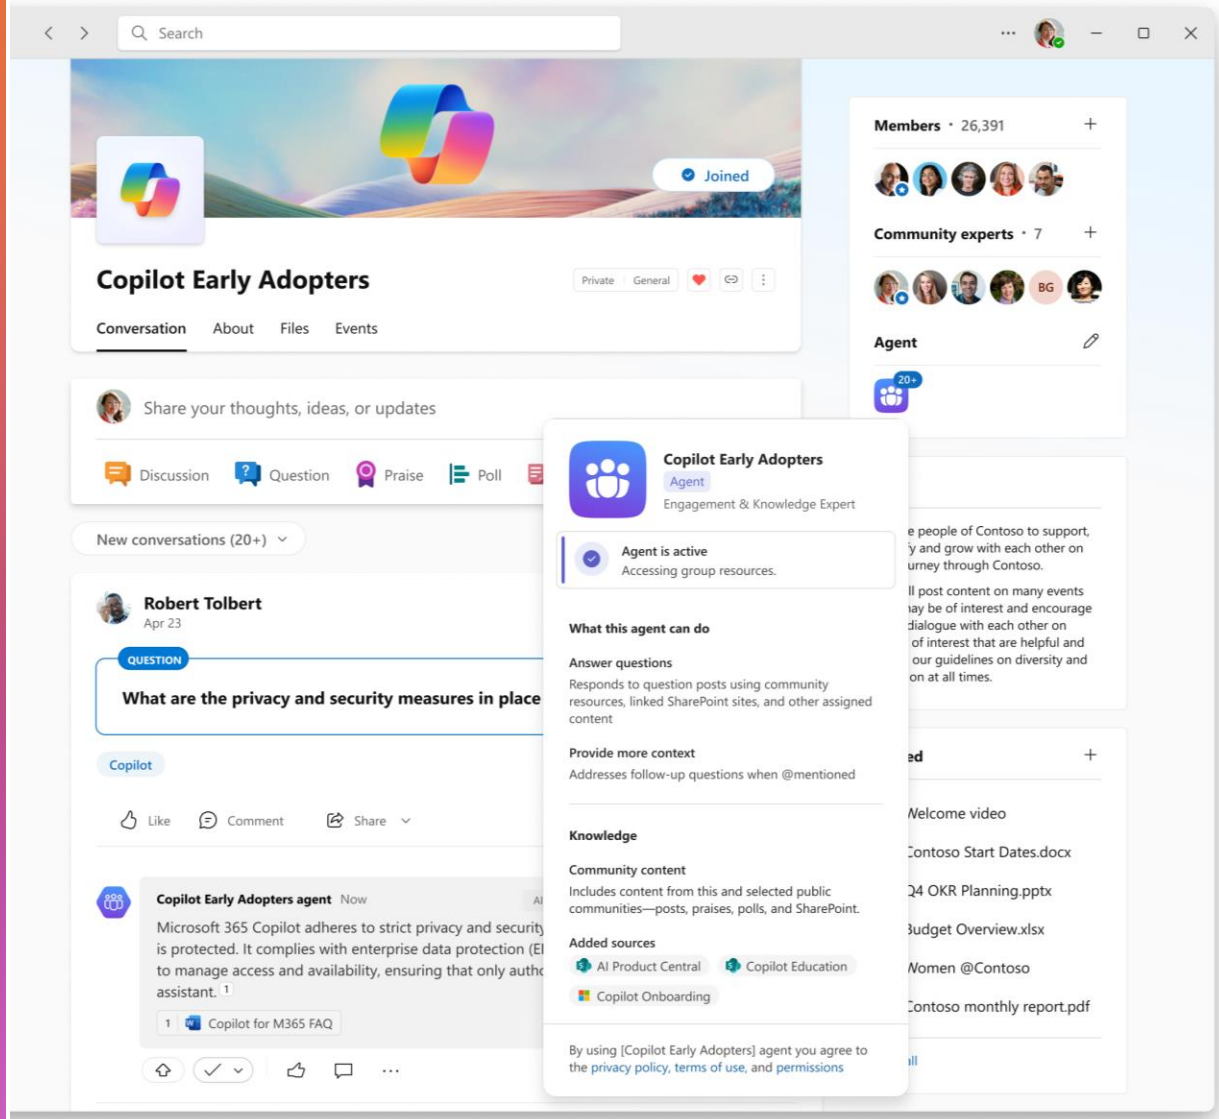

Agenda

- Cycle 5 Scorecard
- Adoption data
- Cycle 6 Priorities
- Special Topics & Risks

Cycle 5 Scorecard

✅ = Complete; ➡️ = Almost ready; ⏳ = In progress; ⭕ = Not started

Goal	Status	Comments
Grounding on community SharePoint	✅	• Rolled out to Public Preview
Grounding on SharePoint outside communities	➡️	• Rolled out to Engage DF; ready for private preview; incl. comms site support
@-mention Private Preview readiness	➡️	• Rolled out to MS Digital today! • FRE coaching bar improvements • Teams notifications for agent responses
Native communities in Teams	➡️	• FE Changes Completed and Bug Bashed. Notifications need to route dynamically (WIP). Nearly ready for FC rollout
Quality: Search Observability Experiments	✅	• ~10% Improvement in LLM relevance and recall against the Eval Set
Quality: A/B Testing Framework in Place	✅	• Short Answer Experiment has been running
Quality: Conciseness and Formatting Experiments	⏳	• Multiple experiments offline evals and surveys. Online experiments ready to go
Quality: Eval Improvements	✅	• Dataset Refinement: Performed a full cleanup of the golden dataset for improved accuracy. • Edge Cases: Discussed and defined a formal process for identifying edge cases and initiated their integration into the golden dataset. • Quality Benchmarking: Increased evaluation thresholds to meet higher quality goals. • Process Optimization: Developed a streamlined method for faster offline evaluations.
Quality: Evaluators	⏳	• Test and integrated coded evaluators for further reading links and specific formatting issues
Rollout what we started in Cycle 4	✅	• Admins & experts can easily see all answers posted by agent • Admins can pause or resume agent; "take over" when owner leaves • 'AI-generated' label and 👍 / 👎 on all agent-posted answers • Squad has detailed product usage dashboard • QoS and Monitoring dashboards, alerting monitors in place
Planning - Agent Identity	✅	• Forward Sync Changes in place. Tech Spec for Engage Impact. AOS POC Planned
Planning - HITL V2, Moderation Gurardrails	✅	• Product Requirements and Designs in place
RAI/DSB	✅	• DSB Approval for Public Preview

Customer Adoption Status

422	1269	200	33.5k	34%
Customers using agents	Enabled communities	Communities w/ agent content	Agent MAU	Auto-posting enabled
+17% MoM	+19% MoM	+11% MoM	+86% MoM	-1 pp MoM

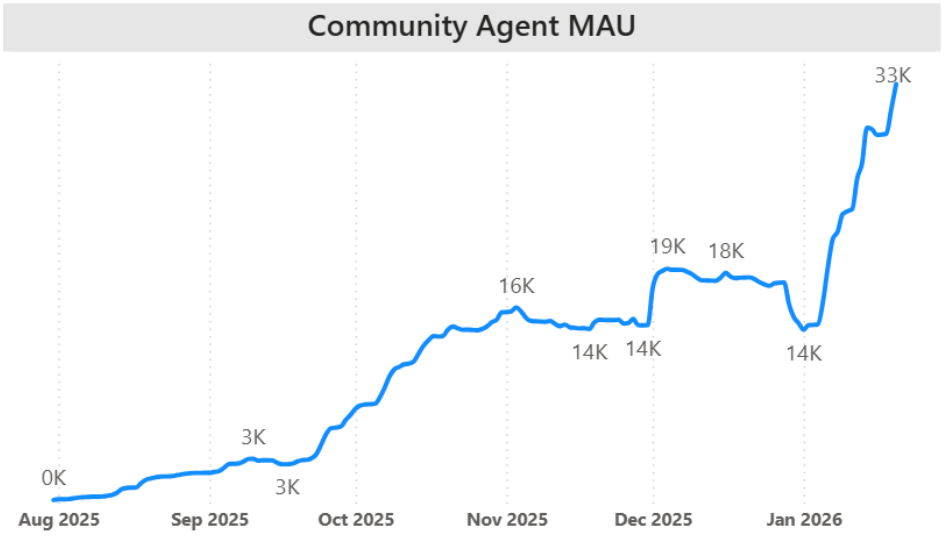

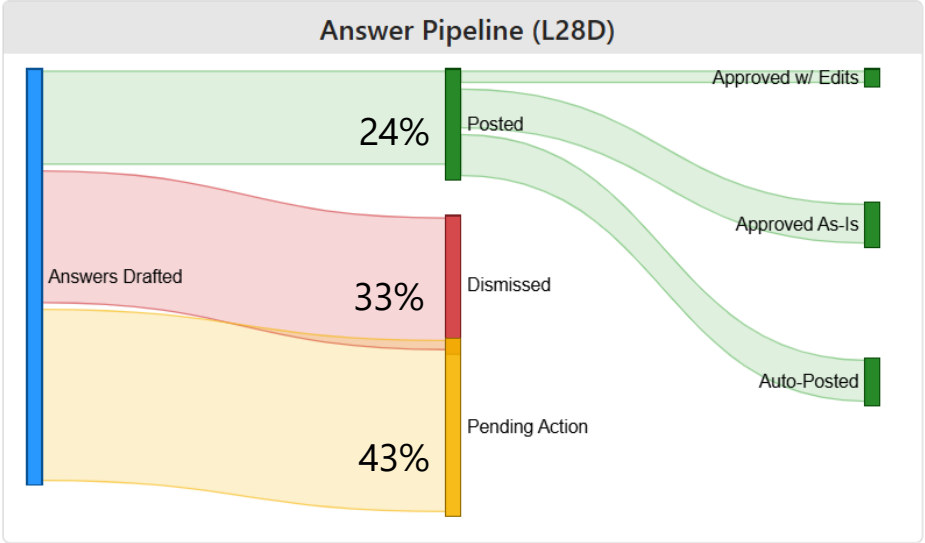

Community Agent Dashboard

Quality Status

- **Refined Golden Dataset and increased thresholds for pass criteria**
  - Fact Retrieval: **69%**
  - Fact Representation: **77%**
  - Cognitive Burden: **49%**
- **Search Improvements**
  - Agent observability logs fine-tuned Search Ranker led to a **+10% increase in relevant posts found**
  - Search experiments all up lead to a **5% increase in questions answered**
- **Online Observability**
  - Reviewer Agent v2 quality insights motivating follow on conciseness/tone experiments
  - Reviewer Agent v2 forecasts:
    - 25% decrease in answer length -> 18% increase in human approval rate
    - 50% decrease in answer length -> 30% increase in human approval rate

CYCLE 6 GOALS

1. Launch agent in Native Communities in Teams
2. Improve answer acceptance rate
   - Answer quality experiments
   - Workflow enhancement experiments
3. Release SharePoint grounding and @-mention to Private Preview
4. Focus on craft, quality, and reliability for GA readiness

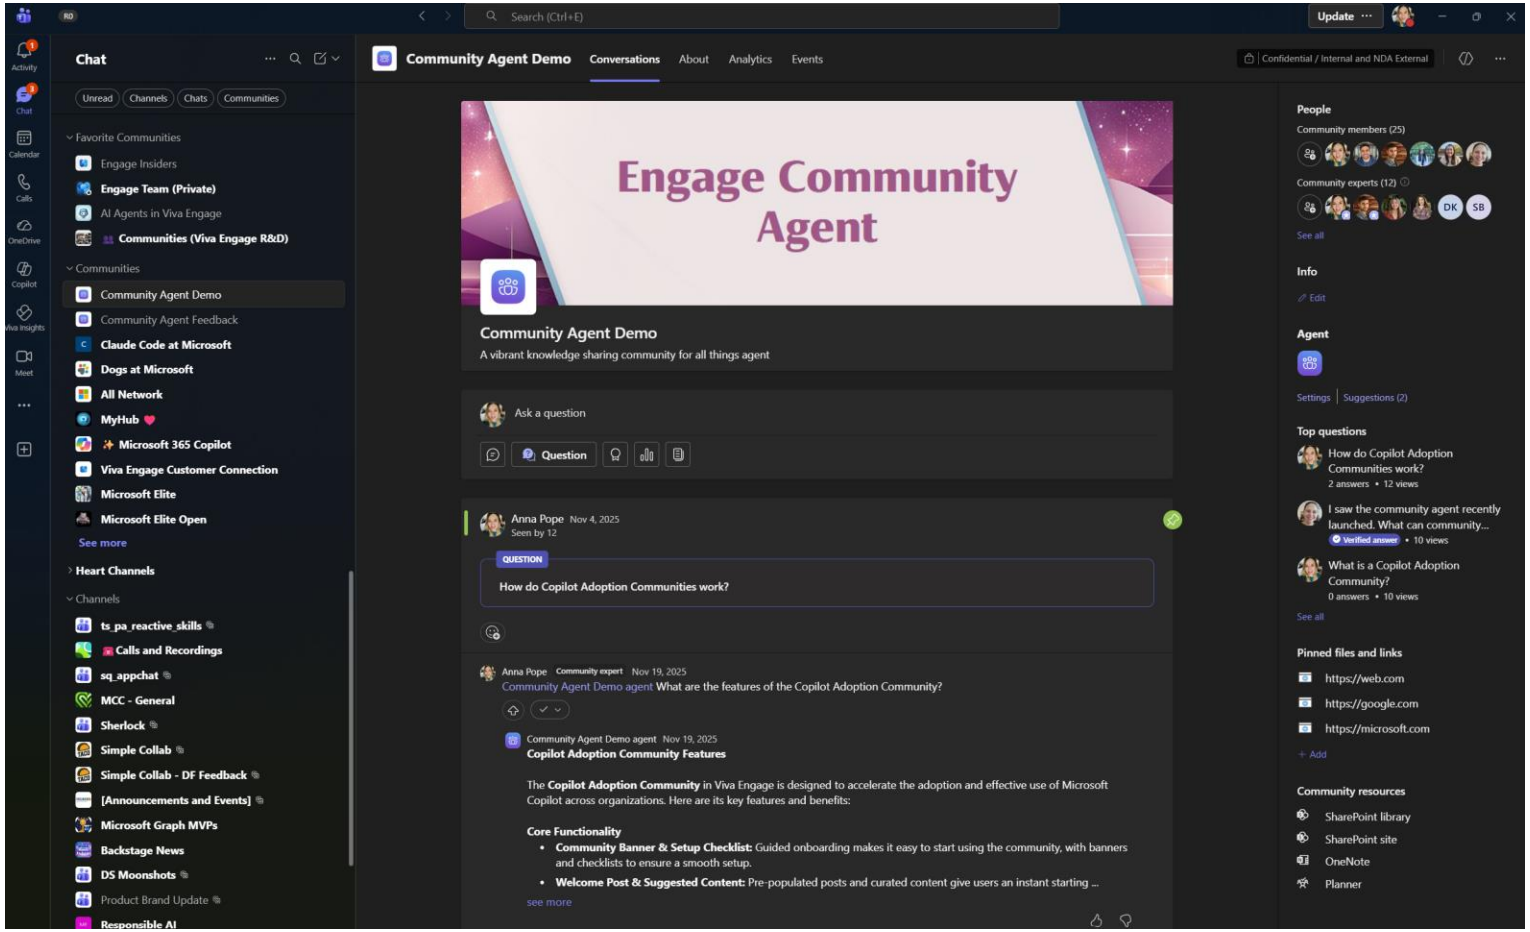

Cycle 6 Release Goals

Feature	Planned Release Ring
Native Communities in Teams integration	Public Preview
@-mention the Community Agent (reactive flow)	Public Preview
Ground on external SharePoint	Public Preview
HITL: Share proposals with reviewers (copy link)	Public Preview
HITL: Remove pending drafts after best/verified answer	Public Preview
HITL: Show queue count and dynamic burndown	Public Preview
Moderation: Skip muted threads	Public Preview

Goal 2: Improve answer acceptance: ‘workflow’ enhancement

Increase reviewer action on agent drafts (approve/edit/dismiss) by making review work easier to notice, faster to complete, and lower-noise.

Experiment/Demo*	Intended Learning
Rank queue items by reviewer relevance 🧪 Prioritize what’s most relevant/important first	Whether ordering review queue by personal relevance increase actions
Auto-remove stale drafts after best/verified Remove redundant proposals when a best/verified answer is marked	Whether removing redundant/stale items makes the queue feel manageable and drives higher action rate
In-feed indicator for pending suggestions 🧪 Surface review work in-context on questions with pending drafts	Whether in-context surfacing increase the % of reviewers who notice a draft exists and take an action
Queue count + dynamic burndown Make progress visible so the queue feels finite; celebrative UI	Whether visible progress reduces the “endless queue” effect and increase persistence / actions per visit
Reviewers can transparently see what questions were skipped or dismissed	—
Lightweight sharing / routing proposals Share with another reviewer so the “right person” acts	Whether explicit routings reduce “not empowered” stalls and increases total actions
One-click edits to reduce reviewer effort** 🧪 E.g. “make concise”, custom instructions	Whether reducing rewrite effort increases approvals and reduces time-to-approve

🧪 *Launch as A/B test*

**Also in our backlog: out-of-UI nudges (Teams, Outlook), motivation & recognition, improvements that make content easier to trust and validate*

***Evaluating offline “pre-experiment” results for conciseness variants to determine if this would be meaningful to pursue*

Goal 2: Improve answer acceptance: answer content & quality

Increase reviewer action by making answer content better—via content experiments, strengthening evaluator-based regression checks and addressing SharePoint grounding issues

Experiment/Item	Intended Learning
Short version (tl;dr first) 🧪 Show a short version of the answer before the rest	Whether a short upfront summary makes answers faster to grok and act on
Conciseness prompt bundle 🧪 Pick best of combine conciseness and tone prompt, integrate conciseness with answer generation prompt	Identify the prompt approach that shortens answers without harming usefulness/accuracy
Use observability signals to suppress low-quality drafts Explore correlations of answer quality signals with approved answers	Whether we can reliably predict approvals and hide low-signal drafts to reduce reviewer load
Model upgrade from 5 → 5.x by Jan 26 deadline	—
Determine extent to which a question is "answered" 🆕 To remove redundant proposals from queue, revise answers accordingly	—
Test and calibrate existing evaluators Cognitive ease, URL integrity, formatting and style	—
Build 2 new evaluators* E.g. Conciseness, tone	—
Investigate and address SharePoint grounding failure modes Irrelevant content, links not clearly supporting answer, docs cited as "Official", broken links or poor anchors	—
Expand eval set Add harder questions, review remaining 30+ datasets	—

**Opportunistic; PMs writing first drafts*

Goal 3: SharePoint Grounding and @-mention in Public Preview

Goal	Details
Deliver grounding on external SharePoint sites to Public Preview	• Support sensitivity labels ✅ • Support Comms sites ✅ • Telemetry • External SharePoint grounding eval approach • SharePoint answer quality improvements
Deliver @-mention reactive flow to Public Preview	• RAI assessment • @-mention content available in Posted view • Request in progress indicator when tagged • Disengage on muted threads

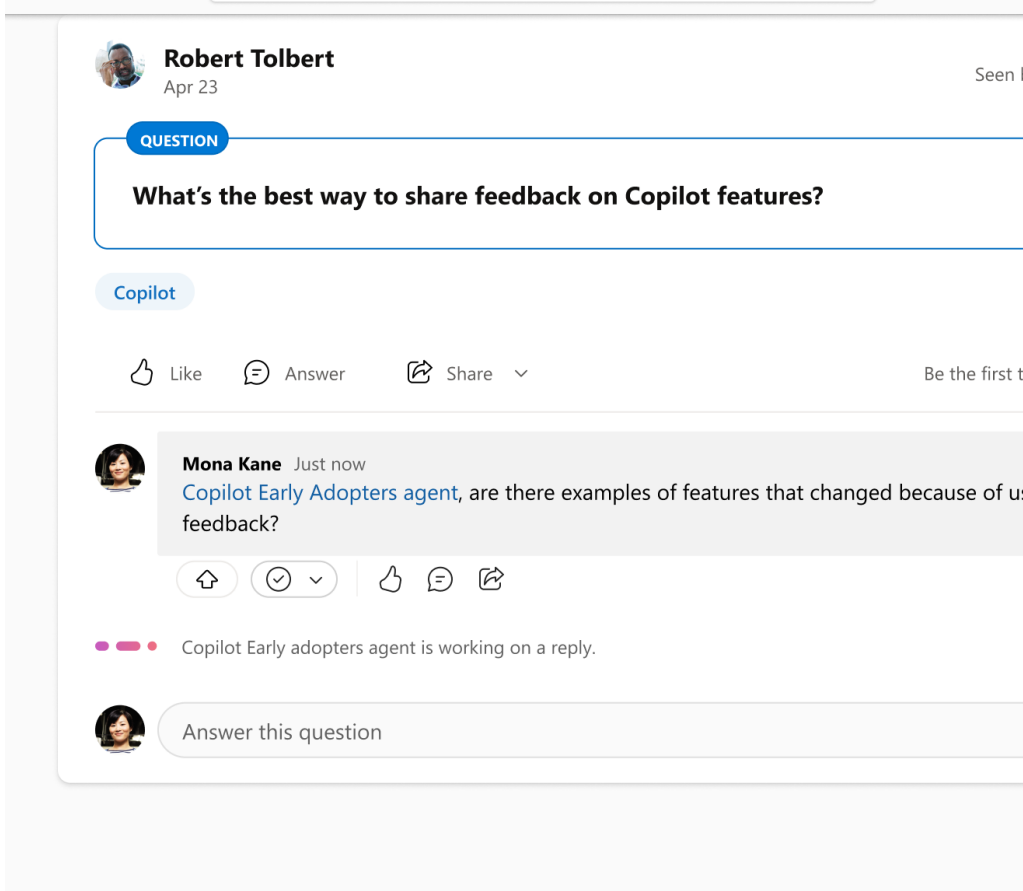

Goal 4: Prepare for GA

Demo/Item
Formalize GA readiness metrics Answer acceptance rate, evaluator metrics, user satisfaction, retention
Harden reliability Ad-hoc, new telemetry and dashboarding
Bring bug backlog in control Keep bug queue within target, address all bugs older than 14 days
Don’t answer ‘muted’ questions
Community can see agent impact in community analytics
Compliance managers can get audit logs including SharePoint grounding events
Support marketing for launch Videos, scripts, comms

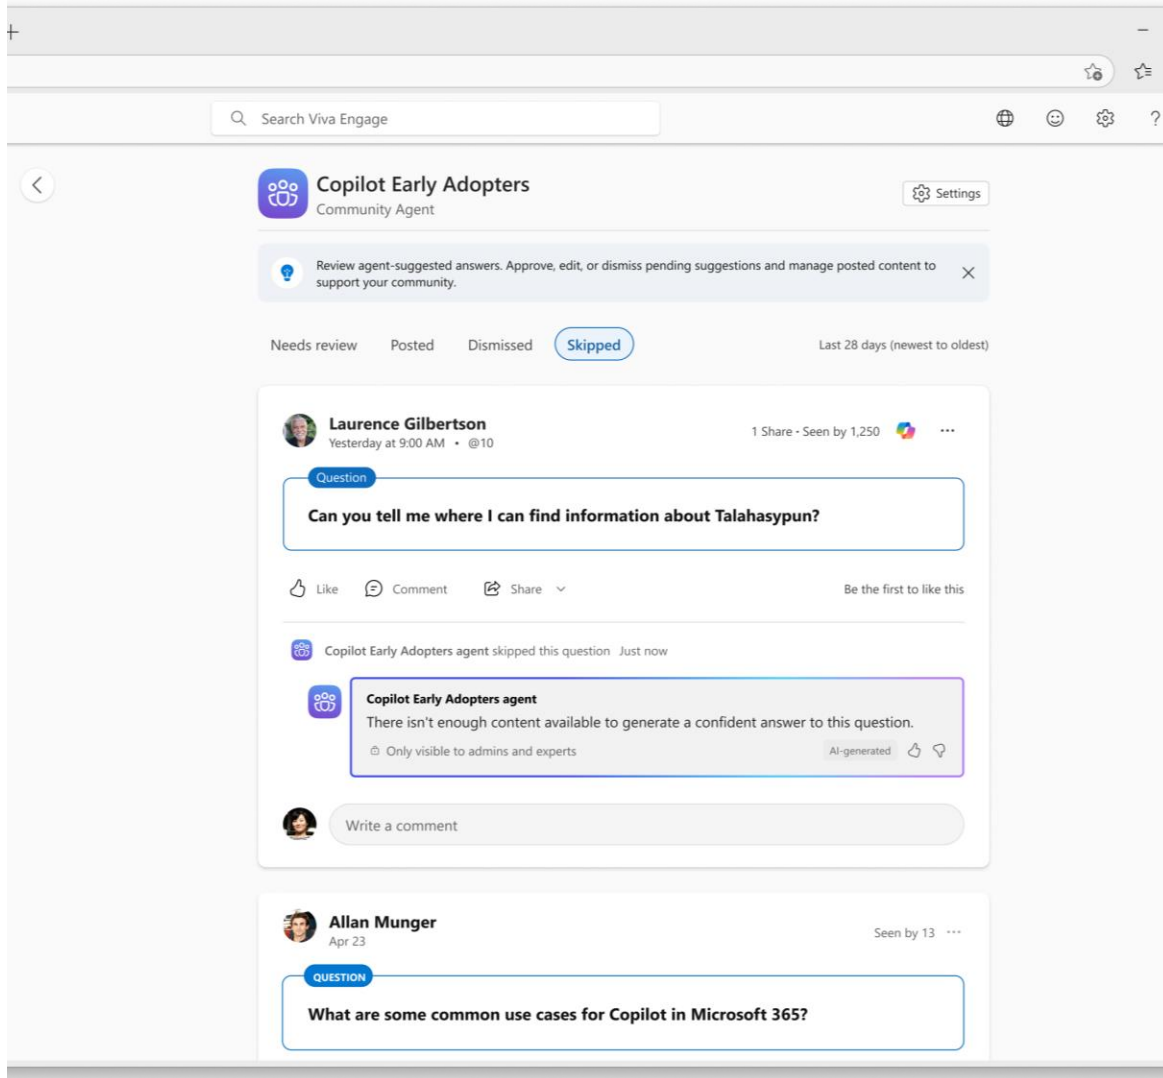

Special Topic: SharePoint Grounding Quality Volatility

- **Enabling community SP grounding has introduced new, SP-specific failure modes**
  1. Irrelevant SP content included
  2. Links not clearly relevant or citing long documents
  3. Link UX issues e.g. broken links, unclear anchor labels
  4. Internal SharePoint documents cited as “official”

- **What we’re doing**
  - Added detailed traces to and evaluating where SP declarative agent issues are occurring
  - Need to then determine which issues can be fixed in our layer vs. require SP DA changes

- **Constraints**
  - Evaluation runs are being constrained by SP DA throttling

Risks

Risk	Mitigation
Reaching statistical significance on quality and workflow experiments can take 4-8 weeks	• Bias for speed with low-risk experiments • Mix in “quick win” experiments with big bets
We are hearing requests from some customers to introduce a tenant level toggle post-GA for change management, compliance reviews, etc.	• Guide customers to seek approvals and test while agent is in preview • *For discussion: introduce a tenant level toggle at GA*
Customer engagement in preview program waning	• Recruit new customers to preview program • Shift to targeted 1:1 calls for deeper feedback
Tight timeline between @-mention Public Preview and GA target date could result in delays if customer feedback signals significant design changes	• If needed, release @-mention to GA after Proactive answering skill
Addressing new failure modes due to SP grounding	• Added detailed tracing + evaluating traces to identify where issues are occurring; this will help us figure out whether SP DA or we need to fix

Quality-based GA readiness

We are quality-driven: GA is primarily gated by measured signals and confidence thresholds

Metric	Why Important	Current	Target
Answer posting rate % of drafted answers that ultimately get posted	Indicates agent is consistently producing publish-worthy content	**23.2%** (+1.8 pts MoM)	**40%**
Answer content quality Measured through offline (and online) evals	Measures presence of failure modes affecting quality as indicated by customers and observed by squad	Fact retrieval: **69%** Fact representation: **77%** Cognitive burden: **49%** +3 more in development	**80%***
Reviewer retention Engaged, acting on drafts	Measures whether admins and reviewers are seeing value, forming habits	**50.9%** (-0.9 pts MoM)	**80%**
👍 **Thumbs up rate** % of thumbs up/down reactions that are positive	Proxy for perceived usefulness by readers	**50.1%** (+7.1 pts MoM)	**70%****

**Aspirational target to pass 4/5 cases; calibrations needed for some plus there will be exceptions like fact retrieval as we adjust conciseness*
***Will materialize as we productionize evaluators*
****Standard across other ODSP agents*